

Sequential Composition for Relaxed Memory

Alan Jeffrey* and James Riely†

*The Servo Project and Roblox

†DePaul University

1. Model

Batty suggest example where dependencies are added and also go away, perhaps by store forwarding. Something like: $(r=x; \ y=1); \ (s=y; \ z=s+r)$

1.1. Preliminaries

The syntax is built from

- a set of *values* \mathcal{V} , ranged over by v, w, ℓ ,
- a set of *registers* \mathcal{R} , ranged over by r, s ,
- a set of *expressions* \mathcal{M} , ranged over by M, N, L .

Memory locations are tagged values, written $[\ell]$. Let \mathcal{X} be the set of memory locations, ranged over by x, y, z .

We require that

- values and registers are disjoint,
- values include at least the constants 0 and 1,
- expressions include at least registers and values,
- expressions do *not* include memory locations.

We model the following language.

$$\begin{aligned} \mu &::= \text{rlx} \mid \text{ra} \mid \text{sc} \\ C, D &::= \text{skip} \mid r := M \mid r := [L]^\mu \mid [L]^\mu := M \\ &\quad \mid \text{fork } G \mid C; D \mid \text{if}(M)\{C\}\text{else}\{D\} \\ G, H &::= 0 \mid \text{thread } C \mid G \parallel H \end{aligned}$$

Memory modes, μ , are relaxed (rlx), release-acquire (ra), and sequentially consistent (sc). Relaxed is the default. *Commands*, C , include reads from and writes to memory at a given mode, as well as the usual structural constructs. *Thread groups*, G , include commands and 0, which denotes inaction. The fork command spawns a thread group. We often drop the words fork and thread.

The semantics is built from the following.

- a set of *actions* \mathcal{A} , ranged over by a ,
- a set of *logical formulae* Φ , ranged over by ϕ, ψ .

We require that

- actions include writes (Wxv) and reads (Rxv),
- formulae include equalities ($M=N$) and ($M=x$),
- formulae are closed under negation, conjunction, disjunction, and substitutions $[M/r]$ and $[M/x]$,
- there is an entailment relation \models between formulae, with the expected semantics.

Logical formulae include equations over locations and registers, such $(x=1)$ and $(r=s+1)$. We use expressions as formulae, coercing M to $M \neq 0$.

Formulae are *open*. Occurrences of register names and memory locations are subject to substitutions of the form $\phi[M/r]$ and $\phi[M/x]$. Actions are not subject to substitution.

We say ϕ *implies* ψ if $\phi \models \psi$. We say ϕ is a *tautology* if $\text{true} \models \phi$. We say ϕ is *unsatisfiable* if $\phi \models \text{false}$.

1.2. Pomsets

We first consider a fragment of our language that can be modeled using simple pomsets.

Definition 1. A *pomset* over \mathcal{A} is a tuple (E, \leq, λ) where

- E is a set of *events*,
- $\leq \subseteq (E \times E)$ is the *causality* partial order,
- $\lambda : E \rightarrow \mathcal{A}$ is a *labeling*.

Let P range over pomsets, and \mathcal{P} over sets of pomsets.

We lift terminology from actions to events. For example, we say that e writes x if $\lambda(e)$ writes x . We also drop quantifiers when clear from context, such as $(\forall e \in E)(\forall x \in \mathcal{X})$.

Definition 2. Action (Wxv) *matches* (Rxw) when $v = w$. Action (Wxv) *blocks* (Rxw) .

We say that e is *fulfilled* if there is a $d \leq e$ which matches it and, for any c which can block e , either $c \leq d$ or $e \leq c$.

We say that P is *fulfilled* if every read in P is fulfilled.

We define *independency* $(\leftrightarrow \subseteq \mathcal{A} \times \mathcal{A})$ as follows.

$$\begin{aligned} \leftrightarrow &= \{(Rxv, Ryw)\} \\ &\cup \{(Wxv, Wyw) \mid x \neq y \vee v = w\} \\ &\cup \{(Rxv, Wyw), (Wxv, Ryw) \mid x \neq y\} \end{aligned}$$

In order to give the semantics, we define several operators over sets of pomsets.

Definition 3.

If $P \in \text{STOP}$ then $E = \emptyset$.

If $P \in (\mathcal{P}_1 \parallel \mathcal{P}_2)$ then $(\exists P_1 \in \mathcal{P}_1) (\exists P_2 \in \mathcal{P}_2)$

- 1) $E = (E_1 \cup E_2)$,
- 2) if $e \in E_1$ then $\lambda(e) = \lambda_1(e)$,
- 3) if $e \in E_2$ then $\lambda(e) = \lambda_2(e)$,
- 4) if $d \leq_1 e$ then $d \leq e$,
- 5) if $d \leq_2 e$ then $d \leq e$,
- 6) E_1 and E_2 are disjoint.

If $P \in (a \rightarrow \mathcal{P})$ then $(\exists P_2 \in \mathcal{P})$

- 1) $E = (E_1 \cup E_2)$,
- 2) if $e \in E_1$ then $\lambda(e) = a$,
- 3) if $e \in E_2$ then $\lambda(e) = \lambda_2(e)$,
- 4) if $d, e \in E_1$ then $d = e$,
- 5) if $d \leq_2 e$ then $d \leq e$,
- 6) if $d \in E_1$ and $e \in E_2$, either $d < e$ or $a \leftrightarrow \lambda_2(e)$.

Using these operators, we can give the semantics for a simple fragment of our language.

$$\begin{aligned} \llbracket 0 \rrbracket &= STOP \\ \llbracket G \parallel H \rrbracket &= \llbracket G \rrbracket \parallel \llbracket H \rrbracket \\ \llbracket x := v; C \rrbracket &= (Wxv) \rightarrow \llbracket C \rrbracket \\ \llbracket r := x; C \rrbracket &= \bigcup_v (Rxv) \rightarrow \llbracket C \rrbracket \end{aligned}$$

If we take $\leftrightarrow = \emptyset$, then we have sequentially consistent execution.

[Do Examples.]

[Do examples with coherence.]

[Note that this allows mumbling for reads and writes.]

[Use refinement (that is subset order) as notion of compiler optimization.]

[Talk about Mazurkiewicz traces.]

1.3. Pomsets with Preconditions

[Problem with previous section is that notion of dependency is impoverished]

Definition 4. A *pomset with preconditions* is a pomset together with $\kappa : E \rightarrow \Phi$.

Definition 5. A pomset with preconditions is *top level* if it is fulfilled and every precondition is a tautology.

Definition 6.

If $P \in STOP$ then $E = \emptyset$.

If $P \in (\mathcal{P}_1 \parallel \mathcal{P}_2)$ then $(\exists P_1 \in \mathcal{P}_1) (\exists P_2 \in \mathcal{P}_2)$

- 1–6) as for \parallel in Definition 3,
- 7) if $e \in E_1$ then $\kappa(e)$ implies $\kappa_1(e)$,
- 8) if $e \in E_2$ then $\kappa(e)$ implies $\kappa_2(e)$.

If $P \in IF(\psi, \mathcal{P}_1, \mathcal{P}_2)$ then $(\exists P_1 \in \mathcal{P}_1) (\exists P_2 \in \mathcal{P}_2)$

- 1–5) as for \parallel in Definition 3 (ignoring disjointness),
- 6) if $e \in E_1 \setminus E_2$ then $\kappa(e)$ implies $\psi \wedge \kappa_1(e)$,
- 7) if $e \in E_2 \setminus E_1$ then $\kappa(e)$ implies $\neg\psi \wedge \kappa_2(e)$,
- 8) if $e \in E_1 \cap E_2$ then $\kappa(e)$ implies $(\psi \wedge \kappa_1(e)) \vee (\neg\psi \wedge \kappa_2(e))$.

If $P \in STOREPRE(x, M, \mathcal{P}_2)$ then $(\exists P_2 \in \mathcal{P}_2) (\exists v \in \mathcal{V})$

- 1–6) as for $(Wxv) \rightarrow P_2$ in Definition 3,
- 7) if $d \in E_1 \setminus E_2$ then $\kappa(d)$ implies $(M=v)$,
- 8) if $e \in E_2$ then either $\kappa(e)$ implies $\kappa_2(e)$ or $e \in E_1$ and $\kappa(e)$ implies $(M=v) \vee \kappa_2(e)$.

If $P \in LOADPRE(x, r, \mathcal{P}_2)$ then $(\exists P_2 \in \mathcal{P}_2) (\exists v \in \mathcal{V})$

- 1–6) as for $(Rxv) \rightarrow P_2$ in Definition 3,

- 7) if $e \in E_2$ then either $e \in E_1$ or $\kappa(e)$ implies $(r=v \vee r=x) \Rightarrow \kappa_2(e)[r/x]$ or $\kappa(e)$ implies $(r=v) \Rightarrow \kappa_2(e)[r/x]$ and $d < e$ for some $d \in E_1$.

Note that when $d \in E_1 \setminus E_2$, *LOADPRE* does not constrain $\kappa(d)$.

The semantics of 0 and \parallel are as before.

$$\begin{aligned} \llbracket \text{if}(\psi)\{C\} \text{ else } \{D\} \rrbracket &= IF(\psi, \llbracket C \rrbracket, \llbracket D \rrbracket) \\ \llbracket r := M; C \rrbracket &= \llbracket C \rrbracket[M/r] \\ \llbracket x := M; C \rrbracket &= STOREPRE(x, M, \llbracket C \rrbracket) \\ \llbracket r := x; C \rrbracket &= LOADPRE(x, r, \llbracket C \rrbracket) \end{aligned}$$

This is essentially the model of Jagadeesan et al. [2020], restricted to relaxed access. There are only two substantial differences. First, for simplicity, we enforce read-read dependencies. Second, in the rule for read prefixing we have substituted $[r/x]$, rather than $[x/r]$. This means that reads clobber local state. We assume registers are only used once—otherwise, one needs to generate a fresh register for the substitution.

With read-read dependencies, the second difference can be seen. For example, the following execution is allowed with $[x/r]$, but not $[r/x]$.

$$x := 0; r := x; \text{if}(r)\{s := x\}; y := s + 1 \parallel x := y$$



[Is there a difference w/o read-read dependencies?]

[Stuff about conditionals and merging events.]

1.4. Pomsets with Predicate Transformers

[The problem with the previous section is that there's no story for sequential composition.]

Definition 7. A *predicate transformer* is a monotone function $\tau : \Phi \rightarrow \Phi$ such that $\tau(\text{false})$ is false, $\tau(\phi \wedge \psi)$ is $\tau(\phi) \wedge \tau(\psi)$, and $\tau(\phi \vee \psi)$ is $\tau(\phi) \vee \tau(\psi)$.

Definition 8. A *family of predicate transformers* indexed by subsets of E consists of predicate transformers τ^D for each set of events D , such that if $C \subseteq D$ then $\tau^C(\phi)$ implies $\tau^D(\phi)$.

Definition 9. A pomset with predicate transformers is a pomset with preconditions, together with a family of predicate transformers τ indexed by subsets of E .

Define *THREAD* to embed pomsets with predicate transformers into pomsets with preconditions simply by dropping the predicate transformer. For the reverse embedding, *FORK* adopts the identity transformer.

Definition 10. If $P \in FORK(\mathcal{P})$ then $(\exists P_1 \in \mathcal{P})$

- 1) $E = E_1$,
- 2) $\lambda(e) = \lambda_1(e)$,
- 3) $\kappa(e)$ implies $\kappa_1(e)$,
- 4) $\tau^D(\phi)$ implies ϕ .

Definition 11. If $P \in STOP$ then $E = \emptyset$ and

- 1) $\tau^D(\phi)$ implies false.

If $P \in SKIP$ then $E = \emptyset$ and

- 1) $\tau^D(\phi)$ implies ϕ .

If $P \in LET(r, M)$ then $E = \emptyset$ and

- 1) $\tau^D(\phi)$ implies $\phi[M/r]$.

If $P \in IF(\psi, P_1, P_2)$ then $(\exists P_1 \in \mathcal{P}_1) (\exists P_2 \in \mathcal{P}_2)$

- 1–8) as for IF in Definition 6,
- 9) $\tau^D(\phi)$ implies $(\psi \wedge \tau_1^D(\phi)) \vee (\neg\psi \wedge \tau_2^D(\phi))$.

If $P \in (P_1 ; P_2)$ then $(\exists P_1 \in \mathcal{P}_1) (\exists P_2 \in \mathcal{P}_2)$,

- 1–5) as for \parallel in Definition 3 (ignoring disjointness),
- 6) if $e \in E_1 \setminus E_2$ then $\kappa(e)$ implies $\kappa_1(e)$,
- 7) if $e \in E_2 \setminus E_1$ then $\kappa(e)$ implies $\kappa'_2(e)$,
- 8) if $e \in E_1 \cap E_2$ then $\kappa(e)$ implies $\kappa_1(e) \vee \kappa'_2(e)$,
where $\kappa'_2(e) = \tau_2^C(\kappa_2(e))$, where $C = \{c \mid c < e\}$,
- 9) $\tau^D(\phi)$ implies $\tau_2^D(\tau_1^D(\phi))$,

If $P \in STORE(x, M)$ then $(\exists v \in \mathcal{V})$

- 1) $\lambda(e) = (Wxv)$,
- 2) $\kappa(e)$ implies $(M=v)$,
- 3) $\tau^\emptyset(\phi)$ implies $\phi[M/x]$,
- 4) $\tau^D(\phi)$ implies $(M=v) \wedge \phi[M/x]$, if $D \neq \emptyset$,
- 5) if $d, e \in E$ then $d = e$.

If $P \in LOAD(x, r)$ then $(\exists v \in \mathcal{V})$

- 1) $\lambda(e) = (Rxv)$,
- 2) $\tau^\emptyset(\phi)$ implies $(r=v \vee r=x) \Rightarrow \phi[r/x]$,
- 3) $\tau^D(\phi)$ implies $(r=v) \Rightarrow \phi[r/x]$, if $D \neq \emptyset$,
- 4) if $d, e \in E$ then $d = e$.

The complete semantics is as follows.

$$\begin{aligned}
\llbracket \text{skip} \rrbracket &= SKIP \\
\llbracket r := x \rrbracket &= LOAD(x, r) \\
\llbracket x := M \rrbracket &= STORE(x, M) \\
\llbracket r := M \rrbracket &= LET(r, M) \\
\llbracket \text{fork } G \rrbracket &= FORK \llbracket G \rrbracket \\
\llbracket C; D \rrbracket &= \llbracket C \rrbracket ; \llbracket D \rrbracket \\
\llbracket \text{if}(\psi)\{C\} \text{ else } \{D\} \rrbracket &= IF(\psi, \llbracket C \rrbracket, \llbracket D \rrbracket) \\
\llbracket 0 \rrbracket &= STOP \\
\llbracket \text{thread } C \rrbracket &= THREAD \llbracket C \rrbracket \\
\llbracket G \parallel H \rrbracket &= \llbracket G \rrbracket \parallel \llbracket H \rrbracket
\end{aligned}$$

[Examples.]

[Skolemization ensures disjunction closure, which is necessary for associativity. Show example.]

2. Complications

[Very drafty. These are just notes.]

- Dependency: preconditions
- TC1: Track local state

2.1. Address Calculation

Definition 12. If $P \in STORE(x, M)$ then $(\exists v \in \mathcal{V})$

- 1) $\lambda(e) = (W[\ell]v)$,
- 2) $\kappa(e)$ implies $(M=v)$,
- 3) $\tau^\emptyset(\phi)$ implies $\phi[M/x]$,
- 4) $\tau^D(\phi)$ implies $(M=v) \wedge \phi[M/x]$, if $D \neq \emptyset$,
- 5) if $d, e \in E$ then $d = e$.

If $P \in LOAD(x, r)$ then $(\exists v \in \mathcal{V})$

- 1) $\lambda(e) = (Rxv)$,
- 2) $\tau^\emptyset(\phi)$ implies $(r=v \vee r=x) \Rightarrow \phi[r/x]$,
- 3) $\tau^D(\phi)$ implies $(r=v) \Rightarrow \phi[r/x]$, if $D \neq \emptyset$,
- 4) if $d, e \in E$ then $d = e$.

2.2. Coherence

Can be encoded in independency, or logic.

- Coherence respects program order: Q_x
- Drop read-read coherence: QW_x (Required for CSE without alias analysis over read only code, not required by hardware)

2.3. Release Acquire

Q

Can be encoded in independency, or logic, but logic is more flexible, and we need that for ARM8.

2.4. ARM Compilation: Internal Acquires


Downgrading acquires/Anton example: $\downarrow x$

2.5. ARM Compilation: Read-read dependencies

RW/RO (control dependencies into reads as in MP with release on right and control dependency on left)

2.6. Redundant Read Elimination

Requires indexing to resolve nondeterminism.



Precondition of $(Wy1)$ is $(r=s)$ in $\llbracket \text{if}(r=s)\{y := 1\} \rrbracket$. Predicate transformers for \emptyset in $\llbracket r := x \rrbracket$ and $\llbracket s := x \rrbracket$ are

$$\begin{aligned}
\langle (r=1 \vee r=x) \Rightarrow \phi[r/x] \mid \phi \rangle, \\
\langle (s=1 \vee s=x) \Rightarrow \phi[s/x] \mid \phi \rangle.
\end{aligned}$$

Combining the transformers, we have

$$\langle (r=1 \vee r=x) \Rightarrow (s=1 \vee s=r) \Rightarrow \phi[s/x] \mid \phi \rangle.$$

Applying this to $(r=s)$, we have

$$\langle (r=1 \vee r=x) \Rightarrow (s=1 \vee s=r) \Rightarrow (r=s) \mid \phi \rangle,$$

which is not a tautology.

Same problem occurs oopsla, where we have:

$$\langle \phi[v/x, r] \wedge \phi[x/r] \mid \phi \rangle,$$

$$\langle \phi[v/x, s] \wedge \phi[x/s] \mid \phi \rangle.$$

Combining the transformers, we have


$$\langle \phi[v/x, r, s] \wedge \phi[v/x, r][x/s] \wedge \phi[x/r][v/x, s] \wedge \phi[x/r, s] \mid \phi \rangle.$$

Applying this to $(r=s)$, we have

$$\langle v=v \wedge v=x \wedge x=v \wedge x=x \mid \phi \rangle,$$

which is not a tautology.

The semantics here allows this by coalescing:



2.7. If Closure

Requires indexing to resolve nondeterminism.

IF closure/case analysis: ψ_e

2.8. Agda

```

κLOAD : Address → Value → Formula
κLOAD a v = RO ∧ Qw[ a ]

τLOADD : Register → Register → Address → Value → Formula → Formula
τLOADD r s a v φ = (value v == register s) ⇒ (φ [ register s / r ] [ register s /[ a ]])

τLOADI : Register → Register → Address → AccessMode → Formula → Formula
τLOADI r s a rlx φ = ¬ Q[ a ] ∧ (RW ⇒ ([ a ] == register s) ⇒ (φ [ register s / r ] [ register s /[ a ]]))
τLOADI r s a ra φ = ¬[ a ] ∧ ¬ Q[ a ] ∧ (RW ⇒ ([ a ] == register s) ⇒ (φ [ register s / r ] [ register s /[ a ]]))

τLOAD∅ : Register → Register → Address → AccessMode → Formula → Formula
τLOAD∅ r s a rlx φ = ¬ Q[ a ] ∧ (φ [ register s / r ] [ register s /[ a ]])
τLOAD∅ r s a ra φ = ¬[ a ] ∧ ¬ Q[ a ] ∧ (φ [ register s / r ] [ register s /[ a ]])

record LOAD (r : Register) (L : Expression) (μ : AccessMode) (P : PomsetWithPredicateTransformers) : Set₁ where

  open PomsetWithPredicateTransformers P

  field a : Event → Address
  field v : Event → Value
  field ψ : Event → Formula

  field d=e : ∀ d e → (d ∈ E) → (e ∈ E) → ((ψ(d) ∧ ψ(e)) ∈ Satisfiable) → (d = e)
  field ℓ=Rav : ∀ e → (e ∈ E) → ℓ(e) = (R (a(e)) (v(e)))
  field κ=κLOAD : ∀ e → (e ∈ E) → (κ(e) = (ψ(e) ∧ (L == address (a(e)) ∧ κLOAD (a(e)) (v(e))))
  field τC=τLOADD : ∀ C φ e → (e ∈ E) → (e ∈ C) → (τ(C)(φ) = (ψ(e) ⇒ τLOADD r (r[ e ] (a(e)) (v(e)) φ))
  field τC=τLOADI : ∀ C φ a e → (e ∈ E) → (e ∉ C) → (τ(C)(φ) = (ψ(e) ⇒ (L == address a) ⇒ τLOADI r (r[ e ] a μ φ))
  field τC=τLOAD∅ : ∀ C φ a s χ → (∀ e → (e ∈ E) → (e ∈ C) → (χ = ¬(ψ(e)))) → (τ(C)(φ) = (χ ⇒ (L == address a) ⇒ τLOAD∅ r s a μ φ))

κSTORE : AccessMode → Expression → Address → Value → Formula
κSTORE rlx M a v = (M == value v) ∧ RW ∧ Q[ a ]
κSTORE ra M a v = (M == value v) ∧ RW ∧ Q

τSTORED : AccessMode → Expression → Address → Value → Formula → Formula
τSTORED rlx M a v φ = (Qw[ a ] ⇒ (M == value v)) ∧ (φ [ M /[ a ] ] [ tt /[ a ] ])
τSTORED ra M a v φ = (Qw[ a ] ⇒ (M == value v)) ∧ (φ [ M /[ a ] ] [ ff /[ a ]* ])

τSTOREI : AccessMode → Expression → Address → Formula → Formula
τSTOREI rlx M a φ = (¬ Qw[ a ]) ∧ (φ [ M /[ a ] ] [ tt /[ a ] ])
τSTOREI ra M a φ = (¬ Qw[ a ]) ∧ (φ [ M /[ a ] ] [ ff /[ a ]* ])

record STORE (L : Expression) (μ : AccessMode) (M : Expression) (P : PomsetWithPredicateTransformers) : Set₁ where

  open PomsetWithPredicateTransformers P

  field a : Event → Address
  field v : Event → Value
  field ψ : Event → Formula

  field d=e : ∀ d e → (d ∈ E) → (e ∈ E) → ((ψ(d) ∧ ψ(e)) ∈ Satisfiable) → (d = e)
  field ℓ=Wav : ∀ e → (e ∈ E) → ℓ(e) = (W (a(e)) (v(e)))
  field κ=κSTORE : ∀ e → (e ∈ E) → (κ(e) = (ψ(e) ∧ (L == address (a(e)) ∧ κSTORE μ M (a(e)) (v(e))))
  field τC=τSTORED : ∀ C φ a e → (e ∈ E) → (e ∈ C) → (τ(C)(φ) = (ψ(e) ⇒ (L == address a) ⇒ τSTORED μ M a (v(e)) φ))
  field τC=τSTOREI : ∀ C φ a χ → (∀ e → (e ∈ E) → (e ∈ C) → (χ = ¬(ψ(e)))) → (τ(C)(φ) = (χ ⇒ (L == address a) ⇒ τSTOREI μ M a φ))

```

Appendix

This paper addresses several errors in [?], which we henceforth refer to as [JJR].

1. Parallel Composition

In [JJR, §2.4], parallel composition is defined allowing coalescing of events. Here we have forbidden coalescing. This difference appears to be arbitrary. In [JJR], however, there is a mistake in the handling of termination actions. The predicates should be joined using \wedge , not \vee .

2. Redundant Read Elimination

In [JJR, §2.6] the semantics of read is defined as follows:

$$\llbracket r := x^{\mu}; C \rrbracket \triangleq \bigcup_v (R^{\mu} x v) \Rightarrow \llbracket C \rrbracket[x/r]$$

The definition of prefixing $((\phi \mid a) \Rightarrow \mathcal{P})$ has several clauses. The most relevant are as follows, where d is the new event labeled with $(\phi \mid a)$ and e is an event from \mathcal{P} :


- (P4C) If d reads v from x then either $e = d$ or $\kappa'(e)$ implies $\kappa(e)[v/x]$.
(P5A) If d reads and e writes then either $\kappa'(e)$ implies $\kappa(e)$ or $d \leq' e$.

We have discovered two issues with this definition.

The first issue concerns the substitution $[x/r]$. It should be $[r/x]$. We noticed this error while developing the alternative characterization presented here. The error causes redundant read elimination to fail in [JJR]. As a result, common subexpression elimination also fails. The problem can be seen in TC2.

$$r := x; s := x; \text{if}(r=s)\{y := 1\} \parallel x := y \quad (\text{TC2})$$


We claimed that TC2 allowed the following execution:



But this execution is not possible using the semantics of [JJR]: (Wy1) has precondition $r=s$ in $\llbracket \text{if}(r=s)\{y:=1\} \rrbracket$. Given the lack of order in the execution, the precondition of (Wy1) must entail $r=1 \wedge r=x$ in $\llbracket s:=x; \text{if}(r=s)\{y:=1\} \rrbracket$. P4C imposes $r=1$, and P5A imposes $r=x$. Adding the second read, the precondition of (Wy1) must entail both $1=1 \wedge 1=x$ and also $x=1 \wedge x=x$. This can be simplified to $x=1$. This leaves a requirement that must be satisfied by a preceding write. Since the preceding write is the initialization to 0, the requirement cannot be satisfied, and the execution is impossible.¹

The substitution $[x/r]$ leaves the obligation on x to be fulfilled by the preceding write. Thus, the read does

1. In [JJR] we ignore the middle terms, mistakenly simplifying this to $1=1 \wedge x=x$. Correcting the error, the attempted execution is:



not update the *value* of x in subsequent predicates. The substitution $[r/x]$, instead, does update the value of x , thus removing any obligation on x for preceding code.

In order to write this, we must update the definition of prefixing reads to include the register. Then P4C becomes:


- (P4C) If d reads v from x then either $e = d$ or $\kappa'(e)$ implies $\kappa(e)[v/r]$.

We can then reason with TC2 as follows: (Wy1) has precondition $r=s$ in $\llbracket \text{if}(r=s)\{y:=1\} \rrbracket$. To avoid introducing order in the execution, the precondition of (Wy1) must entail $r=1 \wedge r=s$ in $\llbracket s:=x; \text{if}(r=s)\{y:=1\} \rrbracket$. P4C imposes $r=1$, and P5A imposes $r=x$. Adding the second read, the precondition of (Wy1) must entail both $1=1 \wedge 1=x$ and also $x=1 \wedge x=x$. This can be simplified to $x=1$. This leaves a requirement that must be satisfied by a preceding write.

With read elimination, the rule for relaxed reads is as follows:

$$\llbracket r := x; C \rrbracket \triangleq \llbracket C \rrbracket[x/r] \cup \bigcup_v (R x v) \Rightarrow_r \llbracket C \rrbracket[r/x]$$


It is interesting to note that the substitution is $[x/r]$ on eliminated reads, and $[r/x]$ on non-eliminated reads. Intuitively, the subsequent value of x is fixed by an explicit read, but not for an eliminated read. In the latter case, the value is fixed by some preceding action. The preceding action may itself be a read. This gives rise to some fear that we might introduce thin-air reads, since we do not enforce read-read coherence. But this is not the case. Consider the following example:



But this is not a problem, since fulfillment requires that (Wx1) precede both reads of x .

3. Internal Acquiring Reads


The second issue concerns acquiring reads. Shortly after publication, Podkopaev [2020] noticed a shortcoming of the implementation on ARM8 in [JJR, §7]. The proof given there assumes that all internal reads can be dropped. However, this is not the case for acquiring reads. For example, [JJR] disallows the following execution, which is allowed by ARM8 and TSO.



The solution we have adopted is to allow an acquiring read to be downgraded to a relaxed read when it is preceded (sequentially) by a relaxed write that could fulfill it. Backporting this solution to [JJR] requires that we add access predicates to the logic and allow

The notion of data-race is incorrect in [JJR].

$x := 1; y^{ra} := 1; r := x^{ra} \parallel \text{if}(y^{ra})\{x^{ra} := 2\}$



Bug is in [Dongol et al., 2019, Lemma A.4]. It assumes that $(R^{ra}x1)$ and $(W^{ra}x2)$ are racing in the first execution because they are not ordered by happens-before. But this is false since neither is plain.

In addition, the ARM8 implementation result given here does not rely on read elimination. Instead we use a recent alternative characterization of ARM8 [Alglave, 2020; Arm Limited, 2020; Alglave et al., 2020].

References

- J. Alglave. This commit adds three alternative formulations of the arm model, both for non-mixed and mixed size accesses. <https://github.com/herd/herdtools7/commit/685ee4b5f821254c947888c6cc731e9eedbe937d>, June 2020.
- J. Alglave, W. Deacon, R. Grisenthwaite, A. Hacquard, and L. Maranget. Armed cats: Formal concurrency modelling at arm. Draft, 2020.
- Arm Limited. Arm architecture reference manual: Armv8, for Armv8-A architecture profile (issue F.c). <https://developer.arm.com/documentation/ddi0487/latest>, July 2020.
- B. Dongol, R. Jagadeesan, and J. Riely. Modular transactions: bounding mixed races in space and time. In J. K. Hollingsworth and I. Keidar, editors, *Proceedings of the 24th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming, PPoPP 2019, Washington, DC, USA, February 16-20, 2019*, pages 82–93. ACM, 2019. doi: 10.1145/3293883.3295708. URL <https://doi.org/10.1145/3293883.3295708>.
- R. Jagadeesan, A. Jeffrey, and J. Riely. Pomsets with preconditions: a simple model of relaxed memory. *Proc. ACM Program. Lang.*, 4(OOPSLA):194:1–194:30, 2020. doi: 10.1145/3428262. URL <https://doi.org/10.1145/3428262>.
- A. Podkopaev. Private correspondence, Nov. 2020.